


calvin and hobbes

by WATKINSON 4973
distributed by UNIVERSAL PRESS syndicate, 6-23



AGENDA

-> Discuss how to avoid 1000s of hours of torture.

Currently in Production

models:/booking-cancellation@xgb-model-depth8

models:/booking-cancellation@xgb-model-star2

Alias!!

models:/booking-cancellation@champion = models:/booking-cancellation@xgb-model-depth8

models:/booking-cancellation@champion = models:/booking-cancellation@xgb-model-star2

Model Registry:

```
mlflow.register_model(f"models:{current_model_id}", PROJECT_NAME) -v1  
mlflow.register_model(f"models:{tuned_model_id}", PROJECT_NAME) - v2
```

```
1 model_server = TestClient(scoring_server.init(champion_model))
2 print("GET /ping →", model_server.get("/ping").status_code)
```

✓ Last Execution: 10:20:12 PM, Duration: 0.0s

inbuilt endpoint health check

GET /ping → 200

time: 8.94 ms (started: 2026-09-30 22:20:12 +05:30)

Generate

+ Code

+ Markdown

```
1 upcoming = pd.read_csv("data/upcoming_arrivals.csv")
2 sample_booking = json.loads(upcoming.loc[upcoming.booking_id == "BK218304", FEATURES].to_json(orient="records"))[0]
3
4 sample_booking # one row per booking, same features and dtypes as train.py uses
```

✓ Last Execution: 10:22:16 PM, Duration: 0.0s


```
response = model_server.post("/invocations", json={"dataframe_records": [sample_booking]})
print(response.status_code, response.json())
```

Post → Endpoint

```
def predict(m)
    a = joblib.load(model)
    return a.predict(m)
```

model.pkl

Test Client (_____)



recommendation data

1. Start an online cloud server (EMR).
2. My server every night at 12:00 AM loads up data from online cloud db (S3), into EMR.
3. Model is trained, the pickle file is dumped into the cloud DB (S3)
4. the model file is loaded up from the S3
5. Deployed on multiple instances.

